

Problem Set 5: Visualizing CNN motifs as PWM

sequence1 = CGATC

A
T
C
G

0	0	1	0	0
0	0	0	1	0
1	0	0	0	1
0	1	0	0	0

(depth=4)

filter =

(3x1)

(depth=1)

0	0	0
0	0	0
0.9	0.1	0
0.1	0.9	0

1.8	0.1	0
-----	-----	---

← apply filter

Sequence2 = ATATCGT

A
T
C
G

1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1

0	0	0	0.1	1.8
---	---	---	-----	-----

← apply filter

max_activation over {seq1, seq2} = 1.8

any ~~activation~~ activation $\geq \frac{1}{2}$ max_activation is "activated"

So "activated" regions from sequence1 = CGA

"activated" regions from sequence2 = CGT

Next, we form alignment
and obtain PWM from
Alignment.

To get PWM, first form frequency table:
Frequencies:

A	0	0	1
G	0	2	0
C	2	0	0
T	0	0	1

⇒

Fractions of each base at each position:

A	0	0	0.5
G	0	1	0
C	1	0	0
T	0	0	0.5

Convert these fractions to "log-odds" against a background model
where each ~~base~~ base occurs w/ probability = $\frac{1}{4}$

Data Augmentation:

- Not enough data $\{(x_i, y_i)\}_{i=1}^N \leftarrow N \leftarrow \text{too small}$
available to get best NN performance

- We know invariants of problem (prior knowledge):
= changes to x_i which should not change y_i

- Idea: Append additional example
 (\tilde{x}_i, y_i) to dataset

where \tilde{x}_i = variant of x_i which we know
should receive same label

- Train NN on larger dataset w/ additional examples

Ex: ① If x_i = sentences, can replace words
with synonyms.

② If x_i = images, can randomly crop/rotate/translate

③ Can sometimes apply dropout to inputs x_i
or add Gaussian noise.

④ (implicit augmentation)

Adversarial Augmentation: Suppose we know y_i should stay the
same if x_i only slightly changes (smooth relationship).

Can use backprop gradients to modify x_i slightly such
that network's prediction changes the most.

Include this (\tilde{x}_i, y_i) in dataset.
Same label as x_i

Transfer Learning:


$D_1 = \{(x_i, y_i)\}_{i=1}^{N_1} \leftarrow \text{Large} = \text{massive data from Task 1}$

Task 2 = main task of interest


$D_2 = \{(x_j, \tilde{y}_j)\}_{j=1}^{N_2} \leftarrow \text{small} = \text{small data from Task 2}$

\uparrow
 x_i, x_j are same features
 y_i, \tilde{y}_j may be different labels!

Step 1: Train Neural Net for Task 1:
 NN output-layer for task 1



Step 2: Replace output layer with output-layer for Task 2



Step 3: Fine-tune the parameters of NN for Task 2.

EX: For classifying biomedical images,
 start with ~~different~~ NN
 trained on Massive object-recognition
 database (ImageNet).

Suppose Task 2 labels not given.

Train Task 2 NN to map $(x_j) \rightarrow (z_j)$
 such that (z_j) and (z_i) (output by Task 1 NN applied to x_i)

follow same statistical probability distribution.

Then apply output layer from Task 1. Match 2 distributions using

Differentiable statistical divergence	
① Maximum Mean Discrepancy	③ Generative Adversarial Network
② Wasserstein Distance	

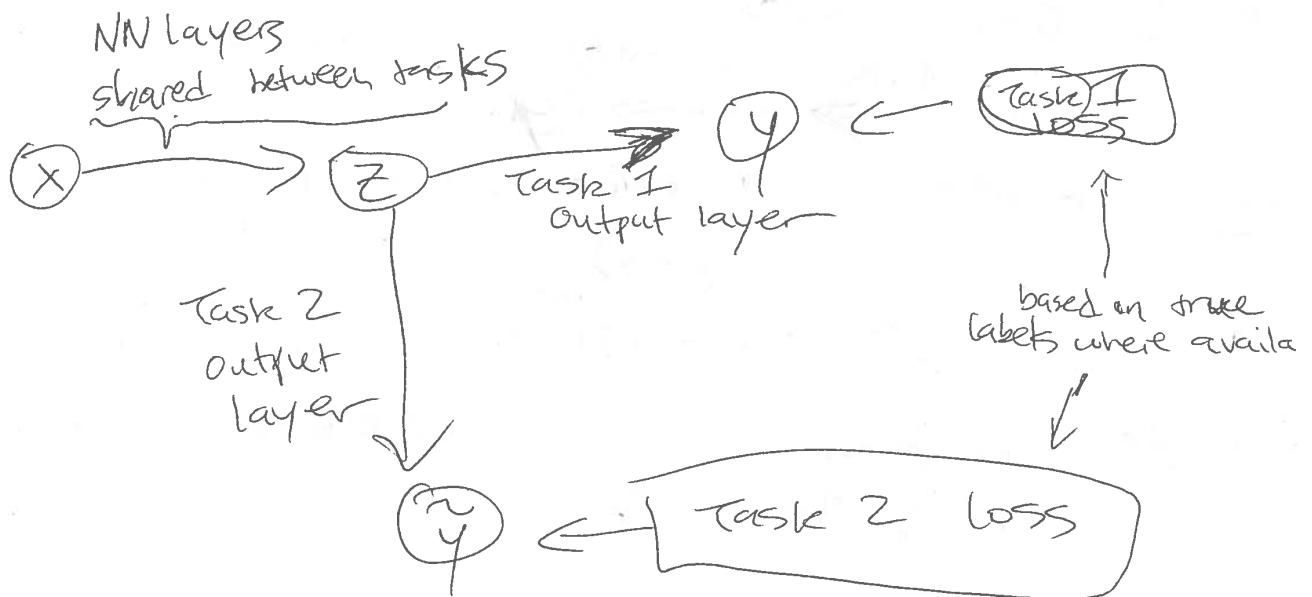
Match 2 distributions using divergence (eq. MMD)

Multi-Task Learning:

2 Tasks: (x, y) , (x, \tilde{y})
Task 1 Task 2

- Not enough data from either task
- We expect similarity between tasks,
how can we "borrow strength" across task

Idea = Train jointly!



When only Task 1 label available, only backpropagate Task 1 loss

When only Task 2 label available, only backpropagate Task 2 loss.

If both labels available, backpropagate both Task 1 & 2 losses.